

STL

01 – Overview and Containers

Contents

STL Contents

Header Files and Naming Conventions

The STD namespace

Aliasing and typedefs

A simple use case

Overview

01 – Overview and Containers

The STL – is a C++ library of

Container classes

Vector, Set, Map, List etc

Algorithms

for_each, transform, find_if

Iterators

Forward, backward,
bi-directional, inserters

Utilities

Adapters – stack, queue,
priority_queue,
Binders, pair, dates & times

Function objects

generators, predicates

Memory Management

New, scoped allocators

Limits

System limits

Error Handling

Exceptions & assertions

String

Narrow, wide, unicode

Numerics

Maths, Random Numbers

Concurrency

Threads, Locks, Tasks, Atomics

A Few Others

Locales, RegEx,

Header Files & Naming convention

Header files in the STL do not have the .h extension

Historically this was the case, STL headers used to have a .h extension, however, these have been deprecated

However, you may still see examples of this in code

In such situations, make every effort to replace these older files with their standard versions

i.e **<vector>** instead of **<vector.h>**
 <iostream> instead of **<iostream.h>** etc

All template classes, template functions in the STL are written in
lower_snake_case

The STD namespace

Every component of the STL
is enclosed in the std namespace
Needs to be scoped to this in order to compile.

There are 3 approaches to this.

1 – Use the entire
namespace
within a scope

```
void UsingEg1()
{
    using namespace std;
    vector<int> ints;
    /// stuff
}
```

2 - Only the whats needed
within a scope

```
void UsingEg2()
{
    using std::vector;
    vector<int> ints;
    /// stuff
}
```

3 - Explicitly scope
as needed

```
void UsingEg3()
{
    std::vector<int> ints;
    /// stuff
}
```

Typedefs

The syntax for templates can become very complicated and difficult to read

To simplify code, use a typedef to give an alias to complicated data types

Use the least amount of scope required

Become familiar with the inbuilt typedefs in the STL

```
void NoTypeDef()
{
    std::vector<int> ints;
    std::vector<int>::const_iterator itr;
    for (itr = ints.begin(); itr != ints.end(); ++itr)
    {
        int item = *itr;
        // Do stuff
    }
}
```

```
void TypedefDemo()
{
    typedef std::vector<int> intColl;
    intColl ints;
    intColl::const_iterator itr;
    for (itr = ints.begin(); itr != ints.end(); ++itr)
    {
        intColl::value_type item = *itr;
        // Do stuff
    }
}
```

Alias Declaration

Modern C++ recommends using an **alias declaration** to create programmer friendly versions of complicated types as opposed to typedefs.

Leads to much simpler code for more advanced C++ programmer and eliminates need for complex use of typename declarations

See E M C++ - Item 9. (*Effective Modern C++* - Scott Meyers)

```
void NoTypeDef()
{
    std::vector<int> ints;
    std::vector<int>::const_iterator itr;
    for (itr = ints.begin(); itr != ints.end(); ++itr)
    {
        int item = *itr;
        // Do stuff
    }
}
```

```
void NamespaceAliasDemo()
{
    using intColl = std::vector<int>;
    intColl ints;
    intColl::const_iterator itr;
    for (itr = ints.begin(); itr != ints.end(); ++itr)
    {
        intColl::value_type item = *itr;
        // Do stuff
    }
}
```


A Simple Use Case

Demonstrate

A container – **vector**

An algorithm – **copy**

An iterator

A utility – **ifstream, ofstream**

Use Case

Build a vector of People objects in memory.

Write the vector to file.

Read the file into a vector

The Header File

```
// STL Headers
#include <vector>

// Application Headers
#include "Person.h"

// Aliases
using People = std::vector<Person>;

// Function Prototypes
void DemoPeople();
void BuildPeople(People& people);
void DumpPeopleToFile(const People& people);
void LoadPeopleFromFile(People& people);
```

Build an in-memory vector

```
void BuildPeople(People& people)
{
    people.clear();

    people.emplace_back("John", "Lennon");
    people.emplace_back("Paul", "McCartney");
    people.emplace_back("George", "Harrison");
    people.emplace_back("Ringo", "Starr");
}
```

Write to File & Read From File

```
void DumpPeopleToFile(const People& people) {  
    std::ofstream store("beatles.txt", std::ofstream::out);  
    for (const auto& person : people)  
        store << person << std::endl;  
    store.close();  
}
```

```
▪ void LoadPeopleFromFile(People& people) {  
    people.clear();  
    std::ifstream store("beatles.txt", std::ios::in);  
  
    std::copy( std::istream_iterator<Person>(store),  
               std::istream_iterator<Person>(),  
               std::back_inserter_iterator<People>(people));  
    store.close();  
}
```

STL Containers

01 – Overview and Containers

Contents

Overview

Sequence Containers & Examples

(Ordered) Associative Containers & Examples

Unordered Associative Containers & Examples

Categories of Containers

The STL containers are divided into one of 4 categories:

Sequence containers

array, vector, deque (pronounced deck), list, forward_list

(Ordered) Associative containers

map, multimap, set, multiset

Unordered Associative containers

unordered_map, unordered_multimap, unordered_set, unordered_multiset

Container adaptors

stack, queue, priority_queue

Some Theory

The STL literature uses the terms **first-class containers** and **near containers**.

A **first class container** is a pure maths / computer science term; when used with the STL it means

Sequence Container AND Ordered Associative AND Unordered Associative

A **near container**, another pure maths / computer science term

When applied to the STL, means other containers that can be used with some or all iterators and can be manipulated using some / all STL algorithms.

Examples of STL near containers
string, bitset, valarray

Container adapters are neither first class nor near containers.

Container Methods

All STL containers have a fairly easy to understand, common and uniform interface.

In general all Containers

- are first class copyable and moveable objects.
- can be constructed with an initial size (pre-allocation)
- automatically re-size when required (with the exception of arrays)

Have methods to generate **const** and **non-const** iterators for moving forward, backwards, random access, inserting and deletion.

- Allow **insertion** – insert an existing object into a container
- Allow **emplacement** – create a brand new object into a container
- Removal

At front, back or in the middle of a container.

Other common methods implemented by most containers

- **size, erase, clear, begin, end, front**

Common typedefs include

- **container::value_type, container::pointer, container::size_type**

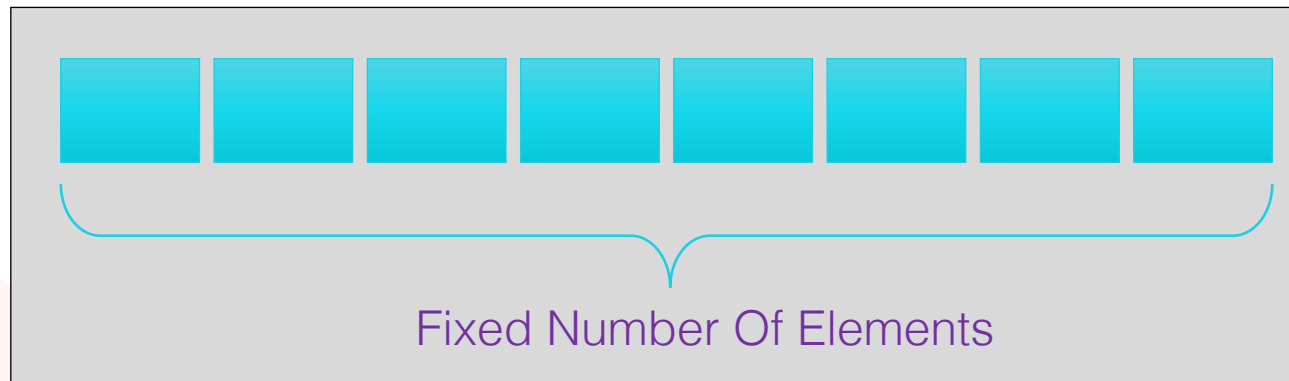
array

Contiguous block of memory

#include <array>

Encapsulates fixed sized arrays

Combines benefits of C-style array (speed, size, etc) with benefits of a standard container (use with algorithms, iterators etc)



array - Example Code

```
std::array<int, 10> arrInt { { 1, 2, 3, 4, 5, 6, 7, 8, 9, 10 } };
```

Constructor here uses
aggregate initialization
Double-braces "{ { " " } }" – Cx11

```
for (const auto& item : arrInt)  
    std::cout << item << ' ' << std::endl;
```

using a range based loop

```
arrInt[0] = 100; arrInt[2] = 100; arrInt[4] = 100; arrInt[6] = 100; arrInt[8] = 100;
```

```
Person john("John", "Lennon");  
Person paul("Paul", "McCartney");  
Person george("George", "Harrison");  
Person ringo("Ringo", "Starr");
```

aggregate initialization
single-braces only (C++14)

```
std::array<Person, 4> arrPeople { john, paul, george, ringo };
```

```
for (const auto& item : arrPeople)  
    std::cout << item << std::endl;
```

using a range based loop

vector


Implemented as an array, a **contiguous** block of elements.

#include <vector>

Allows direct access to any element in the vector

Allows rapid insertions / deletions of items at the back of the vector

Automatically grows when required.



vector- Example Code

```
std::vector<Person> vecBeatles { john, paul, george, ringo };
```

Create a vector

```
std::cout << "Number of Beatles: " << vecBeatles.size() << std::endl;
```

How many elements

```
if (!vecBeatles.empty())
```

Checks if empty

```
    vecBeatles.clear();
```

Clears vector

```
std::vector<Person> anotherVector;
```

```
anotherVector.push_back(john);
```

push_back to add an item to end of vector

```
anotherVector.push_back(paul);
```

```
vecBeatles = anotherVector;
```

assignment

```
if (vecBeatles == anotherVector)
```

equivalence

```
    vecBeatles.push_back(Person("Eric", "Clapton"));
```

```
std::cout << vecBeatles[1] << std::endl;
```

Array-like syntax

```
std::cout << vecBeatles.at(2) << std::endl;
```

At member function

list


Implemented as a **doubly linked list**

```
#include <list>
```

Allows linear access to any element

Compatible with bi-directional iterators

Allows rapid insertions / deletions anywhere in the list



```
std::vector<Person> vecBeatles { john, paul, george, ringo };
```

```
std::list<Person> listBeatles(vecBeatles.begin(), vecBeatles.end());
```

instantiate using the range constructor

```
listBeatles.push_back(Person("Eric", "Clapton"));
```

```
listBeatles.push_front(Person("Bob", "Dylan"));
```

Insert item at front of list

```
std::cout << listBeatles.front() << std::endl;
```

```
listBeatles.pop_front();
```

front and pop_front

```
std::cout << listBeatles.front() << std::endl;
```

```
std::cout << listBeatles.back() << std::endl;
```

```
listBeatles.pop_back();
```

back and pop_back

```
std::cout << listBeatles.back() << std::endl;
```

deque

Best thought of as a **linked list of arrays / vectors**


```
#include <deque>
```

Standard guarantees constant time access to elements


Efficiency is very much dependent on the STL implementation being used.

Allows rapid insertions / deletions of items at the back of the deque

Automatically grows when required.



Internal Structure of a Deque



deque - Example Code

```
std::deque<Person> deqBeatles;
```

```
deqBeatles.emplace_front(john);
```

```
deqBeatles.emplace_front( Person { "Paul", "McCartney" } );
```

Using emplacement to add items to front

```
deqBeatles.emplace_back(george);
```

```
deqBeatles.emplace_back( Person ( "Ringo", "Starr" ) );
```

Using emplacement to add items to back

```
for (const auto& item : deqBeatles)
```

```
    std::cout << item << ' ' << std::endl;
```

```
std::vector<Person> vecBeatles(begin(deqBeatles),end(deqBeatles));
```

using a range based constructor but using the free functions begin and end as opposed to the beginning and end member functions

forward_list


Implemented as a **singly linked list**

```
#include <forward_list>
```

Allows linear access to any element

Allows rapid insertions / deletions anywhere in the list

Provides more space efficient storage as bi-directional iteration not needed



forward_list - Example Code

```
typedef std::forward_list<Person> Beatles;
```

Using a typedef to simplify code later on

```
Beatles beatles;
```

```
// ...
```

```
Beatles::value_type person = beatles.front();
```

Beatles::value_type is a Person

```
std::cout << "First Name = " << person.GetFirstName()  
          << " Last Name = " << person.GetLastName()  
          << std::endl;
```

```
Beatles::pointer pPerson = &beatles.front();
```

```
std::cout << "First Name = " << pPerson->GetFirstName()  
          << " Last Name = " << pPerson->GetLastName()  
          << std::endl;
```

Beatles::pointer is a Person*

set


Implemented as a **balanced red/black binary tree**

#include <set>

Contain a **sorted** set of **unique** objects

No matter how often an item is inserted into a set, the set will only ever contain 1 instance of it.

Objects must be **comparable** (default is to allow < - less than) but comparison operation can be specialised if needed



set - Example Code

```
std::set<Person> setBeatles{ john, paul, george, ringo };

std::cout << "Number of Beatles " << setBeatles.size() << std::endl;

setBeatles.insert(john);
setBeatles.insert(paul);

if (setBeatles.insert(george).second)
    std::cout << "Added " << george << std::endl;
if (!setBeatles.insert(ringo).second)
    std::cout << "Did not add " << ringo << std::endl;

std::cout << "Number of Beatles " << setBeatles.size() << std::endl;
for (const auto& beatle : setBeatles)
    std::cout << beatle << ' ' << std::endl;
```

Add some more without testing for insertion

Add some more with testing for insertion
set::insert returns a **std::pair**
.second member of pair is a Boolean indicating whether insertion was successful or not

Still 4 Beatles

multiset


Implemented as a **balanced red/black binary tree**

#include <set>

Contain a **sorted** set of objects where multiple elements can have equivalent values

Resulting effect of this is that **duplicate items are allowed**

Objects must be **comparable** (default is to allow < - less than) but this can be customized if needed (same as set)



multiset - Example Code

```
std::multiset<Person> multisetBeatles { john, paul, george, ringo };
```

```
std::cout << "Number of Beatles " << multisetBeatles.size() << std::endl;
```

4 Beatles

```
multisetBeatles.insert(john);    multisetBeatles.insert(paul);  
multisetBeatles.insert(george); multisetBeatles.insert(ringo);
```

Add some more Beatles

```
std::cout << "Number of Beatles " << multisetBeatles.size() << std::endl;
```

8 Beatles

```
multisetBeatles.erase(paul);
```

Both Paul McCartneys have been erased

```
std::cout << "Number of Beatles " << multisetBeatles.size() << std::endl;
```

6 Beatles

```
multisetBeatles.clear();
```

0 Beatles

```
std::cout << "Number of Beatles " << multisetBeatles.size() << std::endl;
```

And the Beatles broke up

map

Implemented as a **balanced red/black binary tree**


#include <map>

Contain a **sorted** collection of **key/value** pairs

Maps allow only **unique keys**

The types used for keys and values can differ (more often than not they do differ)

Items used as keys must be **comparable** (default is to allow < - less than) but can be specialised if needed



map - Example Code

```
using BeatleMap = std::map<const std::string, Person>;
```

Alias declaration as opposed to **typedef**

```
BeatleMap mapBeatles;  
// ...
```

How to Add some beatles to a map

```
mapBeatles.insert( { "JL", john } );
```

Using insert

```
mapBeatles.emplace("PMC", paul);
```

Using emplace

```
mapBeatles.insert(BeatleMap::value_type("GH", george));
```

Using **value_type**

```
mapBeatles.insert(std::pair<std::string, Person>("RS", ringo));
```

Using a **pair** directly

```
mapBeatles.insert(std::make_pair("BD", bob));
```

Using a **make_pair**

```
mapBeatles["EC"] = eric;
```

Using array-like syntax

```
std::cout << "Number of Beatles " << mapBeatles.size() << std::endl;
```

```
for (const auto& beatle : mapBeatles)
```

```
    std::cout << beatle.first << " ==> " << beatle.second << std::endl;
```

Print out key ==> value
pairs of entire map.
Beatle is an iterator
First is the key
Second is the value

multimap

Implemented as a **balanced red/black binary tree**


#include <map>

Contain a **sorted** collection of **key/value pairs**

multimaps allow **duplicate** keys

The types used for keys and values can differ (usually the case)

Items used as keys must be **comparable** (default is to allow < - less than) but can be specialised if needed



multimap - Example Code (1 of 2)

```
using BeatlesMultimap = std::multimap<std::string, Person>;
```

```
BeatlesMultimap multimapBealtes;
```

Create a multimap

```
Person john("John", "Lennon");
```

```
Person paul("Paul", "McCartney");
```

```
Person george("George", "Harrison");
```

```
Person ringo("Ringo", "Starr");
```

Create some people

```
Person julian("Julian", "Lennon");
```

```
Person yoko("Yoko", "Ono");
```

```
Person lynda("lynda", "McCartney");
```

```
multimapBealtes.insert(std::make_pair("Lennons", john));
```

```
multimapBealtes.insert(std::make_pair("McCartneys", paul));
```

```
multimapBealtes.insert(std::make_pair("Harrisons", george));
```

```
multimapBealtes.insert(std::make_pair("Starrs", ringo));
```

```
multimapBealtes.insert(std::make_pair("Lennons", julian));
```

```
multimapBealtes.insert(std::make_pair("Lennons", yoko));
```

```
multimapBealtes.insert(std::make_pair("McCartneys", lynda));
```

Add the people to the **multimap**

Note the duplicate keys – multiple **Lennons** and multiple **McCartneys**

multimap - Example Code (2 of 2)

```
std::cout << "Number of Beatles " << multimapBealtes.size() << std::endl;
```

```
// Print out key ==> value pairs of entire map
```

Print out key ==> value pairs of entire map

```
for (const auto& beatle : multimapBealtes)
```

```
    std::cout << beatle.first << " ==> " << beatle.second << std::endl;
```

```
for (auto& itr = multimapBealtes.lower_bound("Lennons");
```

```
    itr != multimapBealtes.upper_bound("Lennons");
```

```
    ++itr)
```

```
    std::cout << itr->first << " ==> "
```

```
        << (*itr).second
```

```
        << std::endl;
```

Just the Lennons

NOTE the **lower_bound** and **upper_bound** member functions.

These return the first and last elements with the key "Lennon"

As iterator in maps and multimaps is a pointer to an **std::pair**, with member variable **first** and **second** It can be used with -> notation OR (*). notation

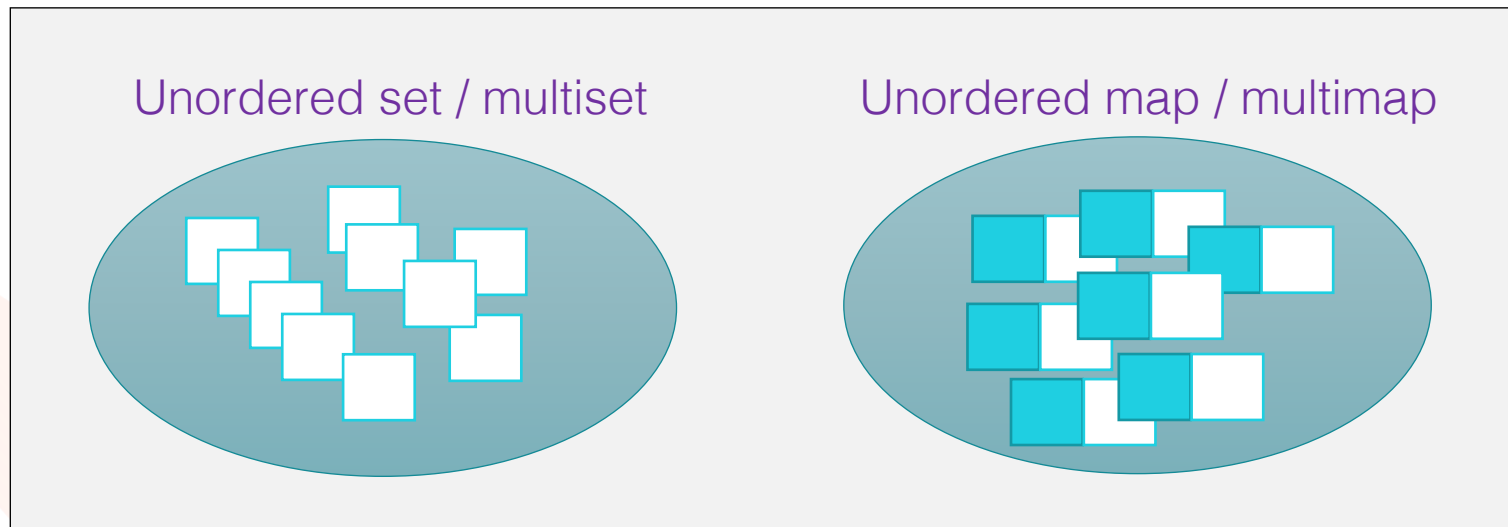
Unordered Containers

The unordered containers in the STL are hash table variants of set/multiset and map/multimap.

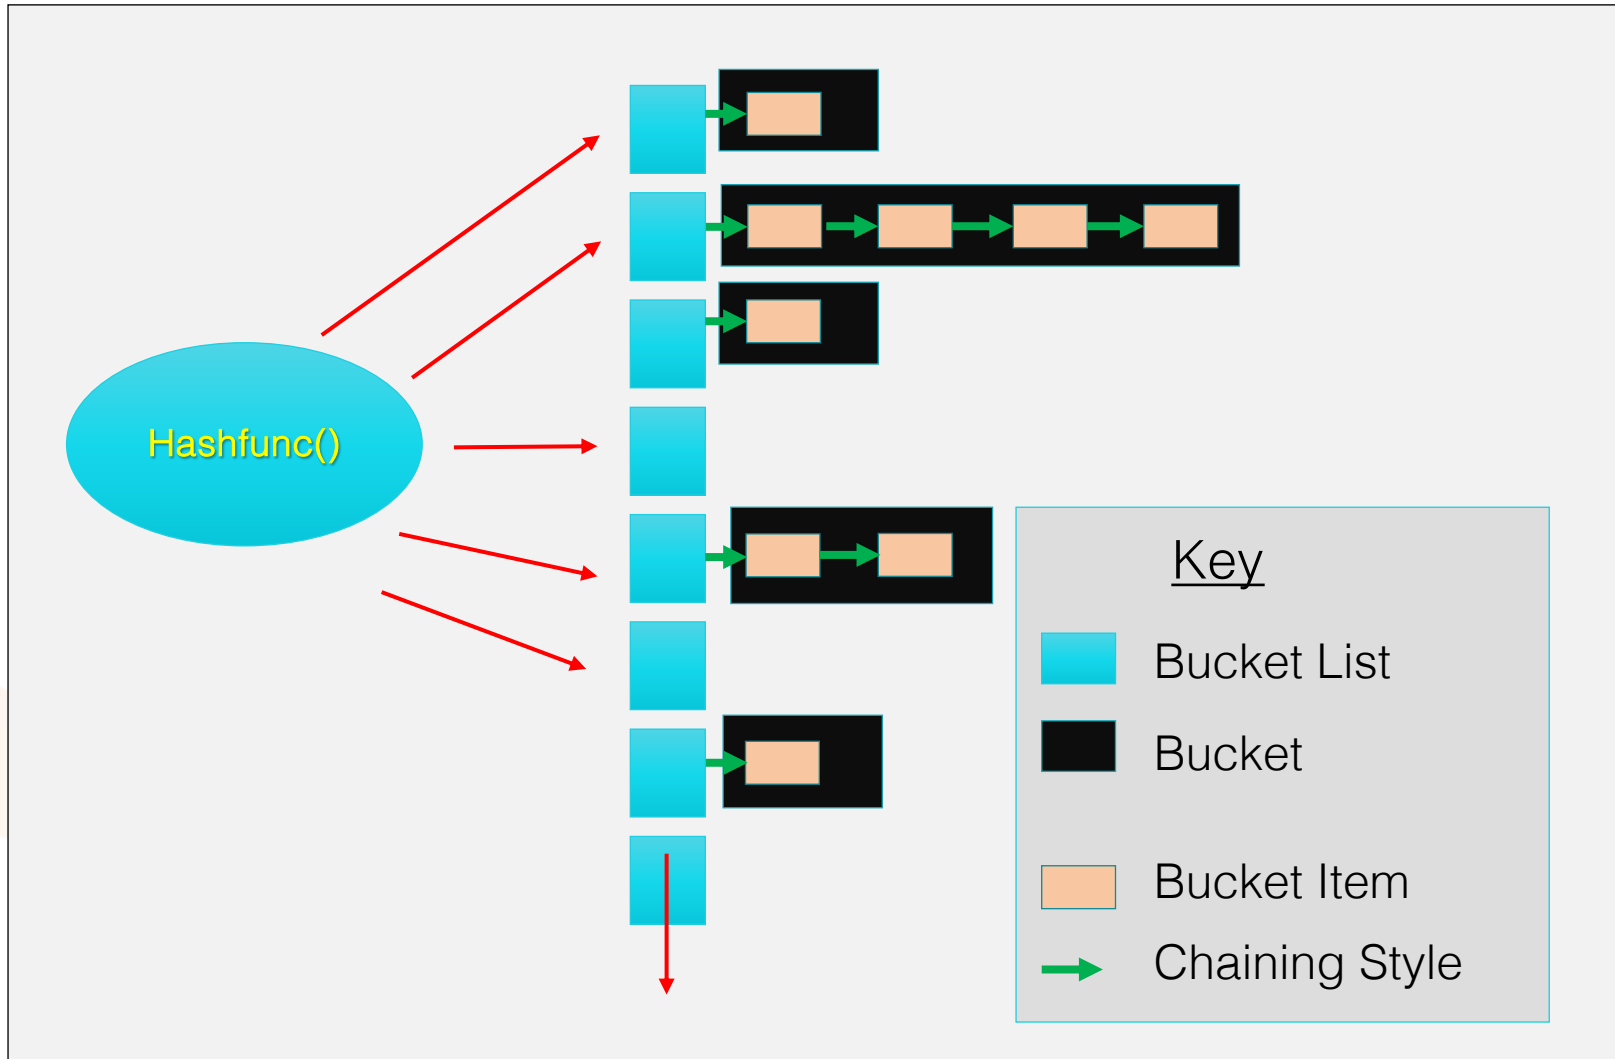
Conceptually, unordered containers contain the same elements as their ordered counterparts except they do not order them.

Unordered sets & multisets store single values, unordered maps and multimaps store key/value pairs.

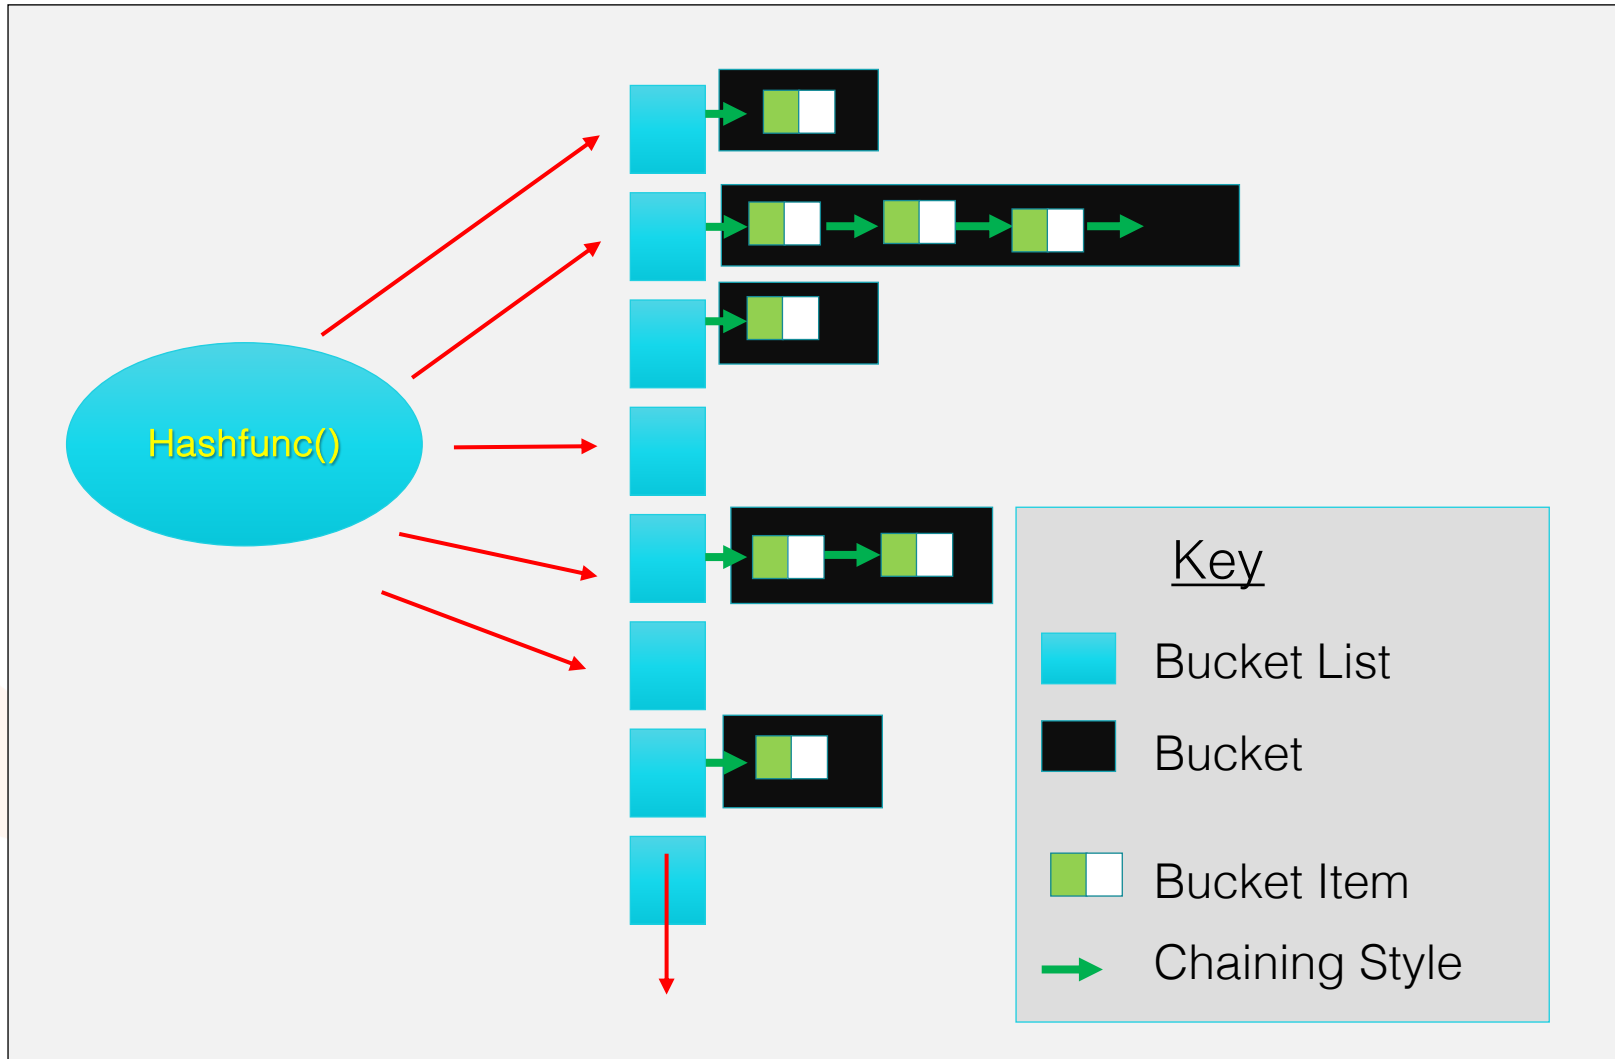
Unordered sets and maps do not allow duplicates, unordered multisets and multimaps do.



Internal Storage of unordered set/multiset



Internal Storage of unordered map/multimap



The Bucket Interface

The unordered containers expose the same sort of interfaces as their ordered counterparts.

In addition, unordered containers expose two further sets of methods

Buckets

bucket_count

Return the number of buckets

max_bucket_count

Return the maximum number of buckets

bucket_size

Return the bucket size

bucket

Locate elements bucket

Hash Policy

load_factor

Return the load factor

max_load_factor

Get or Set maximum load factor

rehash

Set the number of Buckets

reserve

Request a capacity change

Using The Bucket Interface

This example shows the behaviour of

- creating a simple **unordered_map**

- examining its internal structure

- adding some new elements

- examining its internal structure a second time to see the types of changes.

The output shows the internal default hashing function using the following compilers

- VC++ 2015

- LLVM (CLANG) running on Apple

Using The Bucket Interface – Driver Function

```
std::unordered_multimap<std::string, std::string> colours = {
    { "Red", "Richard" },    { "Orange", "Of" },    { "Yellow", "York" },
    { "Green", "Gave" },    { "Blue", "Battle" },  { "Indigo", "In" },
    { "Violet", "Vain" }
};

display_hash_table_state(colours); Initial State

colours.insert( {
    { "Red", "Run" },    { "Orange", "Off" },    { "Yellow", "You" },
    { "Green", "Great" }, { "Blue", "Big" },    { "Indigo", "Irish" },
},
    { "Violet", "Vagabond" }
});

display_hash_table_state(colours); After insertions

colours.max_load_factor(0.7);

display_hash_table_state(colours); After changing load factor
```

Using The Bucket Interface - Helper Templates

```
template <typename K, typename V>
std::ostream& operator << (std::ostream& os, const std::pair<K, V>& p)
{
    return os << "[" << p.first << "," << p.second << "];"
}
```

Using The Bucket Interface - Helper Templates

```
template <typename T>
void display_hash_table_state(const T& container)
{
    std::cout << "size:                " << container.size() <<
    std::endl;
    std::cout << "buckets:                " << container.bucket_count()
    << std::endl;
    std::cout << "load factor:            " << container.load_factor() <<
    std::endl;
    std::cout << "max load factor:       " << container.max_load_factor() <<
    std::endl;

    if ( typeid(typename std::iterator_traits
                <typename T::iterator>::iterator_category) ==
        typeid(std::bidirectional_iterator_tag) )
        std::cout << "chaining style:  doubly-linked" << std::endl;
    else std::cout << "chaining style:singly-linked" << std::endl;

    std::cout << "data:" << std::endl;
    for (auto idx = 0; idx != container.bucket_count(); ++idx) {
        std::cout << "bucket [" << std::setw(2) << idx << "]:";
        for (auto pos = container.begin(idx); pos != container.end(idx); ++pos)
            std::cout << *pos << " ";
        std::cout << std::endl;
    }
}
```

How to determine if
singly linked or
doubly linked

Loop down each bucket

Loop across the elements
of each bucket

The output - Initial State

MS VC++ - Vis Studio 2015

LLVM (on Apple Xcode 6.3)

```
size:          7
buckets:       8
load factor:   0.875
max load factor: 1
chaining style: doubly-linked
data:
    bucket [ 0]:
    bucket [ 1]: [Yellow,York]
    bucket [ 2]: [Violet,Vain]
    bucket [ 3]: [Indigo,In] [Orange,Of]
    bucket [ 4]: [Green,Gave] [Red,Richard]
    bucket [ 5]: [Blue,Battle]
    bucket [ 6]:
    bucket [ 7]:
```

```
size:          7
buckets:       11
load factor:   0.636364
max load factor: 1
chaining style: singly-linked
data:
    bucket [ 0]: [Indigo,In]
    bucket [ 1]: [Green,Gave]
    bucket [ 2]: [Blue,Battle]
    bucket [ 3]:
    bucket [ 4]: [Yellow,York]
    bucket [ 5]:
    bucket [ 6]:
    bucket [ 7]: [Orange,Of]
    bucket [ 8]: [Red,Richard] [Violet,Vain]
    bucket [ 9]:
    bucket [10]:
```

In addition to differences highlights
Note the difference in ordering of items
And in bucket contents

The output - After Insertions

MS VC++ - Vis Studio 2015

```
size:          14
buckets:       64
load factor:   0.21875
max load factor: 1
chaining style: doubly-linked
data:
.....
bucket [11]:   [Orange,Of] [Orange,Off] [Indigo,In]
               [Indigo,Irish]
.....
bucket [18]:   [Violet,Vain] [Violet,Vagabond]
.....
bucket [28]:   [Green,Gave] [Green,Great]
.....
bucket [41]:   [Yellow,York] [Yellow,You]
.....
bucket [45]:   [Blue,Battle] [Blue,Big]
.....
bucket [60]:   [Red,Richard] [Red,Run]
.....
```

LLVM (on Apple Xcode 6.3)

```
size:          14
buckets:       23
load factor:   0.608696
max load factor: 1
chaining style: singly-linked
data:
.....
bucket [2]:    [Red,Richard] [Red,Run]
.....
bucket [4]:    [Blue,Battle] [Blue,Big]
.....
bucket [6]:    [Violet,Vain] [Violet,Vagabond]
.....
bucket [10]:   [Orange,Of] [Orange,Off]
.....
bucket [12]:   [Yellow,York] [Yellow,You]
.....
bucket [20]:   [Indigo,In] [Indigo,Irish]
bucket [21]:   [Green,Gave] [Green,Great]
.....
```

The output - After Changing Load Factor

MS VC++ - Vis Studio 2015

```
size:          14
buckets:       64
load factor:   0.21875
max load factor: 0.7
chaining style: doubly-linked
data:
.....
bucket [11]:   [Orange,Of] [Orange,Off] [Indigo,In]
               [Indigo,Irish]
.....
bucket [18]:   [Violet,Vain] [Violet,Vagabond]
.....
bucket [28]:   [Green,Gave] [Green,Great]
.....
bucket [41]:   [Yellow,York] [Yellow,You]
.....
bucket [45]:   [Blue,Battle] [Blue,Big]
.....
bucket [60]:   [Red,Richard] [Red,Run]
.....
```

LLVM (on Apple Xcode 6.3)

```
size:          14
buckets:       23
load factor:   0.608696
max load factor: 0.7
chaining style: singly-linked
data:
.....
bucket [ 2]:   [Red,Richard] [Red,Run]
.....
bucket [ 4]:   [Blue,Battle] [Blue,Big]
.....
bucket [ 6]:   [Violet,Vain] [Violet,Vagabond]
.....
bucket [10]:   [Orange,Of] [Orange,Off]
.....
bucket [12]:   [Yellow,York] [Yellow,You]
.....
bucket [20]:   [Indigo,In] [Indigo,Irish]
bucket [21]:   [Green,Gave] [Green,Great]
.....
```

